




Token savings come from generated reasoning

Final checkpoints | The same 182 questions solved correctly by both policies



Parsed reasoning is generated text, not a measurement of internal computation. The residual includes tool calls and template framing. All 182 E2 responses are shorter. No jointly correct question uses fewer E2 actions; 148/182 use the same number of decision turns.