



Length shaping also changes how often groups provide a learning signal

Training logs | Trailing 20-update means, shown from update 20 | These are sampled training rollouts



Active group = nonzero composed-reward variation ($1 - \text{frac_reward_zero_std}$). It does not quantify gradient magnitude.
This is not held-out success or an independent replication. The current pair cannot isolate length preference from extra gradient availability.